Creating an LLVM Runtime

EuroLLVM 2026

Michael Kruse

Advanced Micro Devices GmbH

14th April, 2026

Outline

Introduction

Let's Create a Runtime

And They Lived Happily...

AMD
together we advance_

Outline

Introduction
- Why This Tutorial?
- History
- Rationale
- Goal

Let's Create a Runtime

And They Lived Happily...

AMD
together we advance_

Why This Tutorial

[redacted] 1 hour ago Member

After this change, it's no longer possible to run openmp tests in a non-bootstrapping build (i.e. without building clang at the same time).

Previously it was possible to run the tests in a standalone build, by building the `FileCheck` and `not` binaries and pointing at them with `OPENMP_LLVM_TOOLS_DIR` - like I did at mstorsjo/llvm-mingw@ `20260210` /.github/workflows/build.yml#L583-L612.

Now after this change (or before, when doing a runtimes build), running tests only works if there's a `FileCheck` target already in the build, i.e. if doing a bootstrapping build:

llvm-project/openmp/cmake/OpenMPTesting.cmake
Lines 52 to 74 in c6425aa

```
52    if (${OPENMP_STANDALONE_BUILD})
53      find_standalone_test_dependencies()
54
55      # Set lit arguments.
56      set(DEFAULT_LIT_ARGS "-sv --show-unsupported --show-xfail")
57      if (MSVC OR XCODE)
58        set(DEFAULT_LIT_ARGS "${DEFAULT_LIT_ARGS} --no-progress-bar")
59      endif()
60      if ("${CMAKE_SYSTEM_NAME}" MATCHES "AIX")
61        set(DEFAULT_LIT_ARGS "${DEFAULT_LIT_ARGS} --time-tests --timeout=3000")
62      endif()
63      set(OPENMP_LIT_ARGS "${DEFAULT_LIT_ARGS}" CACHE STRING "Options for lit.")
```


AMD
together we advance_

Why This Tutorial

[redacted]

[redacted] on Dec 8, 2025 • edited

Yep, just for context from the Rust side. We have a lot of builders in CI, but the slowest one we have is the x86 dist builder, which we use to e.g. create compiler artifacts. We already build LLVM there around 4 times due to PGO and Bolt, so we really try hard not to make it any slower, since we are already at ~3.5 hrs, and we run this CI before any merge into main.

One of the compromises there is that we have two checkouts of LLVM - once as a standalone llvm/llvm-project build (A) and once as a submodule of rustc in `src/tools/llvm-project`. We use build A to compile C/C++ code, and keep rebuilding the second checkout (B) during bootstrapping.
We only build clang as part of our LLVM Build A. However, rustc is based on LLVM B, and as such we'd also want the openmp and offload libraries to also be part of our LLVM Build B, which does not build the in-tree clang to save compile times.

We've tried for quite a bit to get the runtimes to build despite no in-tree clang or lld in LLVM B, and failed, I'm happy to share all the issues I ran into there. A standalone build instead allows me to use clang from build A to build offload+openmp from checkout B, that's compatible with what we have in rustc. So until that is fixed, standalone builds are likely the only builds that our infra team would be ok with.

AMD
together we advance_

Why This Tutorial

[redacted] on Dec 8, 2025 • edited

Yep, just for context from the Rust side. We have a lot of builders in CI, but the slowest one we have is the x86 dist builder, which we use to e.g. create compiler artifacts. We already build LLVM there around 4 times due to PGO and Bolt, so we really try hard not to make it any slower, since we are already at ~3.5 hrs, and we run this CI before any merge into main

[redacted] 1 hour ago • edited Contributor

We were using the OpenMP standalone build mode in order to build the static libomp.a runtime library which cannot be build together with the shared libomp.so (there is no FLANG_RT_ENABLE_SHARED and FLANG_RT_ENABLE_STATIC equivalent for libomp, the LIBOMP_ENABLE_SHARED flag is binary: one can get either shared or static library, but never both in the same build).

The OpenMP standalone build mode has been removed by llvm/llvm-project#182022

This complicates our situation. Considering almost nonexistent demand for the static OpenMP runtime library, it is not worth pursuing it.

AMD
together we advance_

Why This Tutorial

[redacted] on Dec 8, 2025 • edited

Yep, just for context from the Rust side. We have a lot of builders in CI, but the slowest one we have is the x86 dist builder, which we use to e.g. create compiler artifacts. We already build LLVM there around 4 times due to PGO and Bolt, so we really try hard not to make it any slower, since we are already at ~3.5 hrs and ...

[redacted] on Jan 7 • edited

Member

That comes with two drawbacks: 1) checking out `llvm` alone (I'm not talking about `llvm-project`) takes a lot of time. 2) the way to pass CMake arguments to runtime build is extremely user unfriendly. If these two can't be improved, I strongly object the removal of standalone build.

We we[redacted] build t[redacted] for libc[redacted] build).

The Op[redacted]

This co[redacted] our situation. Considering almost nonexistent demand for the static OpenMP runtime library, it is not worth pursuing it.

[redacted]om

AMD
together we advance_

Why This Tutorial

[redacted] on Dec 8, 2025 • edited

Yep, just for context from the R

[redacted] on Dec 8, 2025

Member

Rust colleagues rely on standalone builds for their offloading extensions (see rust-lang/rust#148671).

CC @ZuseZ4

the removal of standalone build.

build).

The Op

This co

Considering almost nonexistent demand for the static OpenMP runtime library, it is not worth pursuing it.

time. 2) the

be improved, I strongly object

om

AMD
together we advance_

Why This Tutorial

[redacted] on Dec 8, 2025 • edited

Yep, just for context from the [redacted]

[redacted] on Jul 26, 2025

It is definitely not the same!
For standalone build, I used to not needing the llvm-project tree. I could build libomp by just downloading the OpenMP release tarball: `llvmorg-21.1.0-rc1` openmp-21.1.0-rc1.src.tar.xz (download)

I just realized that this workflow got bricked by #62240

Member

Member

Member

time. 2) the

will be improved, I strongly object

omplicates our situation. Considering almost nonexistent demand for the static OpenMP runtime library, it is not worth pursuing it.

om

AMD
together we advance_

Why This Tutorial

[redacted] on Dec 8, 2025 • edited

Yep, just for context from the [redacted]

Member

[redacted] on Jan 7

Member

I don't have a build tree. That's the whole point of building standalone. And the installed binary tree doesn't work.

I just realized that this w[redacted]

[redacted] complicates our situation. Considering almost nonexistent demand for the static Openm[redacted]

pursuing it.

AMD
together we advance_

History

1. Initially, each subproject had its own SVN repository
   - Subprojects categorized as "tools" (clang: 2009), "runtime" (libprofile), and "projects" (e.g. llvm-gcc)
   - An example project boilerplate existed

[1] https://reviews.llvm.org/D20992
[2] https://lists.llvm.org/pipermail/llvm-dev/2020-October/145997.html
[3] https://reviews.llvm.org/D26365
[4] https://reviews.llvm.org/D132478

AMD
together we advance_

History

1. Initially, each subproject had its own SVN repository
   - Subprojects categorized as "tools" (clang: 2009), "runtime" (libprofile), and "projects" (e.g. llvm-gcc)
   - An example project boilerplate existed
2. "runtime" subdirectory disappears with dysfunct PGO removal (2013)

[1] https://reviews.llvm.org/D20992
[2] https://lists.llvm.org/pipermail/llvm-dev/2020-October/145997.html
[3] https://reviews.llvm.org/D26365
[4] https://reviews.llvm.org/D132478

AMD
together we advance_

History

1. Initially, each subproject had its own SVN repository
   - Subprojects categorized as "tools" (clang: 2009), "runtime" (libprofile), and "projects" (e.g. llvm-gcc)
   - An example project boilerplate existed
2. "runtime" subdirectory disappears with dysfunct PGO removal (2013)
3. Desire to separate runtimes[1,2] (2016–2020)
   - Somewhat different from compile-time utils/tools/compilers

[1] https://reviews.llvm.org/D20992
[2] https://lists.llvm.org/pipermail/llvm-dev/2020-October/145997.html
[3] https://reviews.llvm.org/D26365
[4] https://reviews.llvm.org/D132478

AMD
together we advance_

History

1. Initially, each subproject had its own SVN repository
   - Subprojects categorized as "tools" (clang: 2009), "runtime" (libprofile), and "projects" (e.g. llvm-gcc)
   - An example project boilerplate existed
2. "runtime" subdirectory disappears with dysfunct PGO removal (2013)
3. Desire to separate runtimes[1,2] (2016–2020)
   - Somewhat different from compile-time utils/tools/compilers
4. Introduction of `LLVM_ENABLE_PROJECTS`[3] (2016)
5. Introduction of `LLVM_ENABLE_RUNTIMES`[4] (2022)
   - First for libc++

[1] https://reviews.llvm.org/D20992
[2] https://lists.llvm.org/pipermail/llvm-dev/2020-October/145997.html
[3] https://reviews.llvm.org/D26365
[4] https://reviews.llvm.org/D132478

AMD
together we advance_

Rationale

Why LLVM_ENABLE_RUNTIMES?

Build Host

The system we build on

- cmake, ninja, cc, …
- llvm-tblgen, mlir-tblgen, …

Cross-Compile Target

Where Clang will run

- clang, opt, lld, …
- LLVMSupport{.a/.so}, LLVMPolly.so, …

Compiler Target

What Clang will compile for

- a.out
- compiler-rt, libc, libc++, …

- Automate multiple build-steps
  - Instead of: Build clang → build compiler-rt for x86_64 → build compiler-rt for arm64 → …
- Reduce per-runtime boilerplate/differences between build modes
- Runtimes compiled (Stage 2) using just-built Clang

AMD
together we advance_

Rationale

Why LLVM_ENABLE_RUNTIMES?

Build Host

The system we build on

- cmake, ninja, cc, …
- llvm-tblgen, mlir-tblgen, …

Cross-Compile Target

Where Clang will run

- clang, opt, lld, …
- LLVMSupport{.a/.so}, LLVMPolly.so, …

Compiler Target

What Clang will compile for

- a.out
- compiler-rt, libc, libc++, …

Stage 1

Stage 2

- Automate multiple build-steps
  - Instead of: Build clang → build compiler-rt for x86_64 → build compiler-rt for arm64 → …
- Reduce per-runtime boilerplate/differences between build modes
- Runtimes compiled (Stage 2) using just-built Clang

AMD
together we advance_

Tutorial Runtime

What is our tutorial runtime going to do?

Our Runtime

```
void _dublin_hello_world() {
  puts("Dia daoibh, a dhomhain!\n");
}
```

- Use LLVM best-practices
- Linux only
- Prioritize conciseness over
  - ... features
  - ... comments
  - ... boilerplate not topic in this tutorial
- Point out shortcomings of the current system

AMD
together we advance_

Outline

Introduction

Let's Create a Runtime
- Setup
- Step 1: Register with the LLVM Build System
- Step 2: Build a Library Artifact
- Step 3: An Example that uses the Library
- Step 4: Build Modes
- Step 5: CMake Cache Files
- Step 6: Artifact Output Location
- Step 7: Installation
- Step 8: Shared and Static Libraries
- Step 9: Regression Tests
- Step 10: Unittests
- Step 11: Sphinx Docs
- Step 12: Doxygen Docs
- Step 13: Cross-Compilation
- Step 14: Accelerator Offloading
- Step 15: Depending on Other LLVM Libraries / C++
- Step ?

And They Lived Happily...

AMD
together we advance_

Setup

https://github.com/Meinersbur/llvm-dublin/tree/dublin-0-main

- Update Readme for tutorial
  - `README.md`
- Make git ignore `install*` subdirs (like `build*` already is)
  - `.gitignore`

```
alias cmake='cmake -G Ninja -DCMAKE_C_COMPILER_LAUNCHER=ccache -DCMAKE_CXX
```


AMD
together we advance_

Register with the LLVM Build System

https://github.com/Meinersbur/llvm-dublin/tree/dublin-1-basic

- Allow `LLVM_ENABLE_RUNTIME=dublin`
  - `llvm/CMakeLists.txt`
  - `runtimes/CMakeLists.txt`
- Subproject boilerplate
  - `dublin/CMakeLists.txt`
  - `dublin/CODE_OWNERS.TXT`
  - `dublin/LICENSE.TXT`

AMD
together we advance_

Step 2: Build a Library Artifact

- Library source files
  - `lib/Dublin/dublin.c`
  - `include/dublin/Dublin/dublin.h`
- Build `libdublin.a`
  - `add_executable(dublin dublin.c)`

AMD
together we advance_

Standard Subproject Layout

```
subproject
├─ docs
│  └─ *.rst, *.md
├─ examples
│  └─ Example1
│     └─ *.c, *.cpp
├─ test
│  ├─ Unit
│  │  └─ lit.cfg.py, lit.site.cfg.py.in
│  ├─ Component1
│  │  └─ *
│  ├─ Component2
│  │  └─ *
│  └─ lit.cfg.py, lit.site.cfg.py.in
├─ unittests
│  ├─ Component1
│  │  └─ Component1Test.cpp
│  └─ Component2
│     └─ Component1Test.cpp
├─ cmake
│  ├─ Modules/modules
│  │  └─ *.cmake
│  └─ caches
│     └─ *.cmake
├─ include
│  └─ subproject
│     ├─ Component1
│     │  └─ component1.h
│     └─ Component2
│        └─ component2.h
├─ lib
│  ├─ Component1
│  │  ├─ component1.c
│  │  └─ component1-internal.h
│  └─ Component2
│     └─ component2.c
└─ utils
   └─ utility
      └─ utility.cpp
```


AMD
together we advance_

An Example that uses the Library

https://github.com/Meinersbur/llvm-dublin/tree/dublin-3-example

- Example source
  - examples/dublin-hello/dublin-hello.c
- Build instructions
  - examples/dublin-hello/CMakeLists.txt

AMD
together we advance_

Build Modes

- Detecting Legacy Build Modes
  - `CMakeLists.txt`

AMD
together we advance_

Build Modes

Legacy Build Modes

As project: `cmake -S <llvm-project>/llvm -D LLVM_ENABLE_PROJECTS=dublin`
Standalone: `cmake -S <llvm-project>/dublin`

- Intended to be superseded by runtimes build modes[1]
- If your build infrastructure still uses this, please consider updating

Runtimes Build Modes

Bootstrapping: `cmake -S <llvm-project>/llvm -D LLVM_ENABLE_RUNTIMES=dublin`
Default/Standalone: `cmake -S <llvm-project>/runtimes -D LLVM_ENABLE_RUNTIMES=dublin`

- Shared+common boilerplate code in `<llvm-project>/runtimes/CMakeLists.txt`

[1] https://lists.llvm.org/pipermail/llvm-dev/2020-October/145997.html

AMD
together we advance_

Bootstrapping Runtimes Build Mode

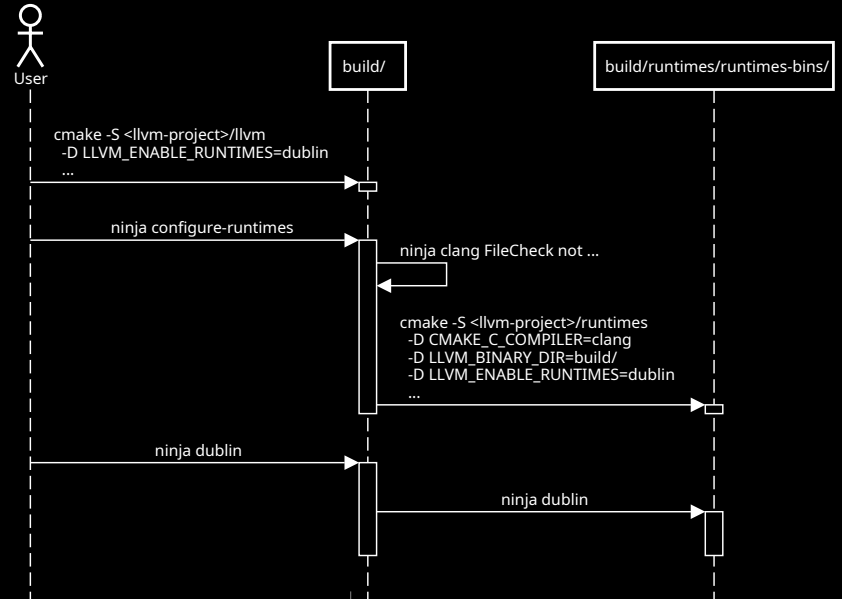

AMD
together we advance_

Default/Standalone Runtimes Build Mode

Skip the Middleman

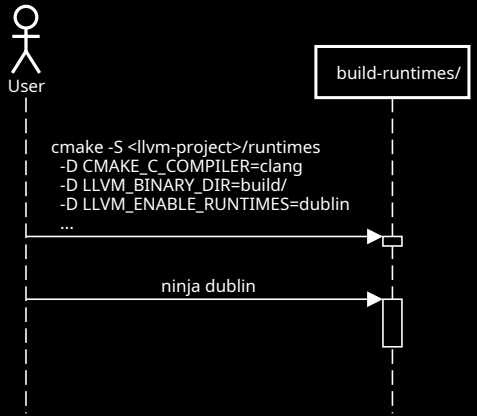

AMD
together we advance_

Default/Standalone Runtimes Build Mode

runtimes/CMakeLists.txt

```
project(Runtimes C CXX ASM)

find_package(LLVM HINT ${LLVM_BINARY_DIR})
# more boilerplate code

foreach (runtime ${LLVM_ENABLE_RUNTIMES})
  add_subdirectory(${runtime})
endforeach ()
```


AMD
together we advance_

Bootstrapping Build Mode

llvm/runtimes/CMakeLists.txt

- Configure nested builddirs
  - runtimes/buitins-bins for Compiler-RT
  - runtimes/runtimes-bins
- Forward configuration parameters
  - Selected configuration options (e.g. LLVM_INCLUDE_TESTS)
  - Any option that starts with the capitalized runtime name (i.e. DUBLIN_*)
  - The content of RUNTIMES_CMAKE_ARGS
    e.g. -DRUNTIMES_CMAKE_ARGS="*-DCMAKE_CXX_FLAGS=-ferror-limit=1*"
- Build prerequisites
  - Test infrastructure: llvm-lit, FileCheck, not, …
  - Compilers: Clang, Flang (if enabled via CMAKE_ENABLE_PROJECTS)
- Build targets for the nested builddir(s)
  - Control targets: configure-runtimes, install-runtimes, clean-runtimes, …
  - Added dependency: all, install
  - Just forwarded: dublin, check-dublin, runtimes → all

AMD
together we advance_

CMake Cache Files

https://github.com/Meinersbur/llvm-dublin/tree/dublin-5-cachefiles

- `cmake/caches/runtimes-bins.cmake`
  - Pre-populated CMake parameters as a bootstrapping build would set them

AMD
together we advance_

Artifact Output Location

https://github.com/Meinersbur/llvm-dublin/tree/dublin-6-location

- `cmake/Modules/GetToolchainDirs.cmake`
  - Determine standard locations
- `lib/Dublin/CMakeLists.txt`
  - Set artifact output location

AMD
together we advance_

Installation

▶

AMD
together we advance_

Shared and Static Libraries

- Add a Configuration option
-

AMD
together we advance_

Regression Tests

▶

AMD
together we advance_

Unittests

▶

AMD
together we advance_

Sphinx Docs

▶

AMD
together we advance_

Doxygen Docs

▶

AMD
together we advance_

Cross-Compilation

▶

AMD
together we advance_

Accelerator Offloading

▶

AMD
together we advance_

Depending on Other LLVM Libraries / C++

▶

AMD
together we advance_

Platform Introspection

▶

AMD
together we advance_

Outline

Introduction

Let's Create a Runtime

And They Lived Happily...

More Topics

Build System Improvements

Conclusion

AMD
together we advance_

Uncovered Topics

- Platform Support: Windows/MacOS/AIX/Fuchsia/...
- Multilib
- ...

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric
- Writing runtimes in C++

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric
- Writing runtimes in C++
- compiler-rt builtins special handling

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric
- Writing runtimes in C++
- compiler-rt builtins special handling
- GPUs like any other target

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric
- Writing runtimes in C++
- compiler-rt builtins special handling
- GPUs like any other target
- Fat libraries

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric
- Writing runtimes in C++
- compiler-rt builtins special handling
- GPUs like any other target
- Fat libraries
- Get rid of `LLVM_ENABLE_PER_TARGET_RUNTIME_DIR`

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric
- Writing runtimes in C++
- compiler-rt builtins special handling
- GPUs like any other target
- Fat libraries
- Get rid of `LLVM_ENABLE_PER_TARGET_RUNTIME_DIR`
- QEMU (`CMAKE_CROSSCOMPILING_EMULATOR`)

AMD
together we advance_

Future Improvements

- Inform LLVM developers and distributions
- `LLVM_ENABLE_RUNTIMES` Documentation
- Propagate `RUNTIMES_*` options
- Unify the CMake code of all the runtimes
  - Use `libdublin` as a template?
- Make Less libcxx-centric
- Writing runtimes in C++
- compiler-rt builtins special handling
- GPUs like any other target
- Fat libraries
- Get rid of `LLVM_ENABLE_PER_TARGET_RUNTIME_DIR`
- QEMU (`CMAKE_CROSSCOMPILING_EMULATOR`)
- ...

AMD
together we advance_

Conclusion

- Now you know how `LLVM_ENABLE_RUNTIMES` works
- Use the Knowledge Wisely
- I currently cannot recommend working on improving the system, waste of your time

AMD
together we advance_

AMD